



# Suneeta Mall

Rambling of a curious engineer & data scientist

[Posts](#) · [Projects](#) · [Talks](#) · [About](#)

# End-to-end reproducible Machine Learning pipelines on Kubernetes

Monday, December 23, 2019, 12:00 AM [Machine-learning](#), [AI](#), [Kubernetes](#), [Reproducible-ml](#)

This is [Part 3](#) - **End-to-end reproducible Machine Learning pipelines on Kubernetes** of technical blog series titled [Reproducibility in Machine Learning](#). [Part 1](#) & [Part 2](#) can be found [here](#) & [here](#) respectively.

Change Anything Changes Everything (CAKE) principle -[Scully et al](#) is real in ML. Also, 100% reproducible ML code comes at a cost of speed - a non-negotiable aspect in today's time. If we cannot be 100% and change is evident, then the only way to maintaining explainability, understanding, trust & confidence is through version control everything.



Figure 1: Version control explained by XKCD

In this post, we will be looking at building an end to end fully automated ML pipeline that can maintain full provenance across entire ML system. In my view, a standard machine learning workflow looks like one below:

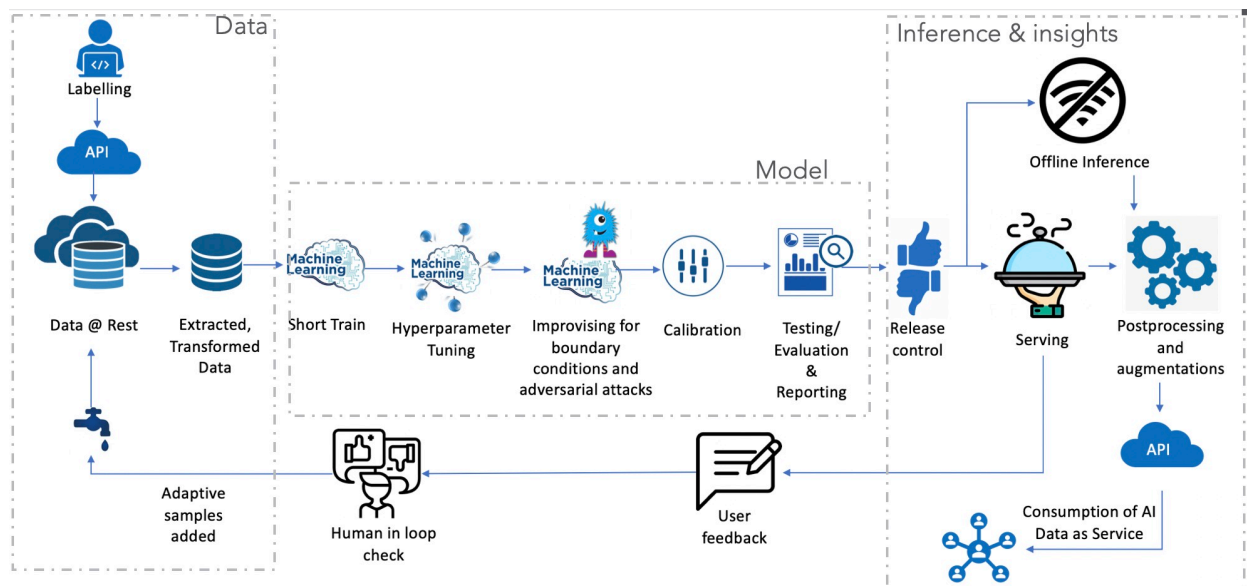


Figure 2: Machine Learning end to end system

So we will be working towards building this system with full provenance over it. For this, we will be extending [our sample semantic segmentation example](#) based on [Oxford Pet dataset](#)

To build this ML workflow, we will be using [Kubernetes](#) - a container orchestration platform. On top of [Kubernetes](#), we will be running [Pachyderm](#) a software that will do the heavy lifting of maintaining provenance across data, artifacts and ml processes.

## What to version control

In [part 1](#) of this blog series, we discussed above the challenges, shown in figure 3, in realizing reproducible ML.

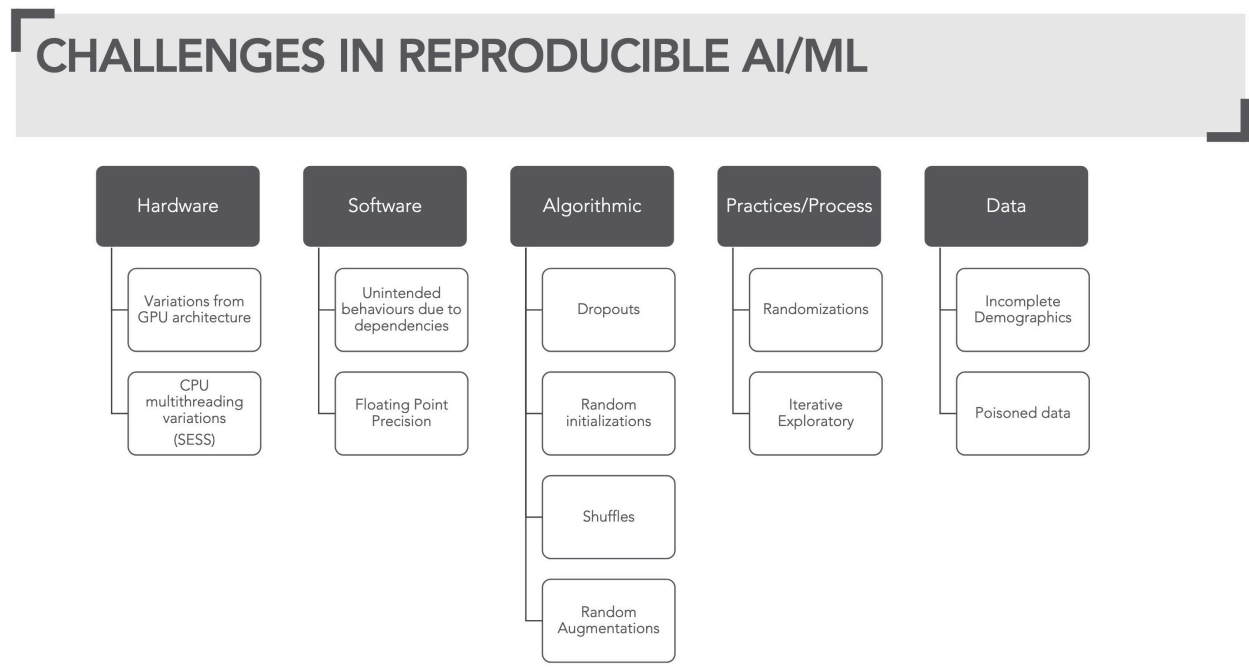


Figure 3: Overview of challenges in reproducible ML

Presence of these challenges in system wide view of ML is shown in figure 4.

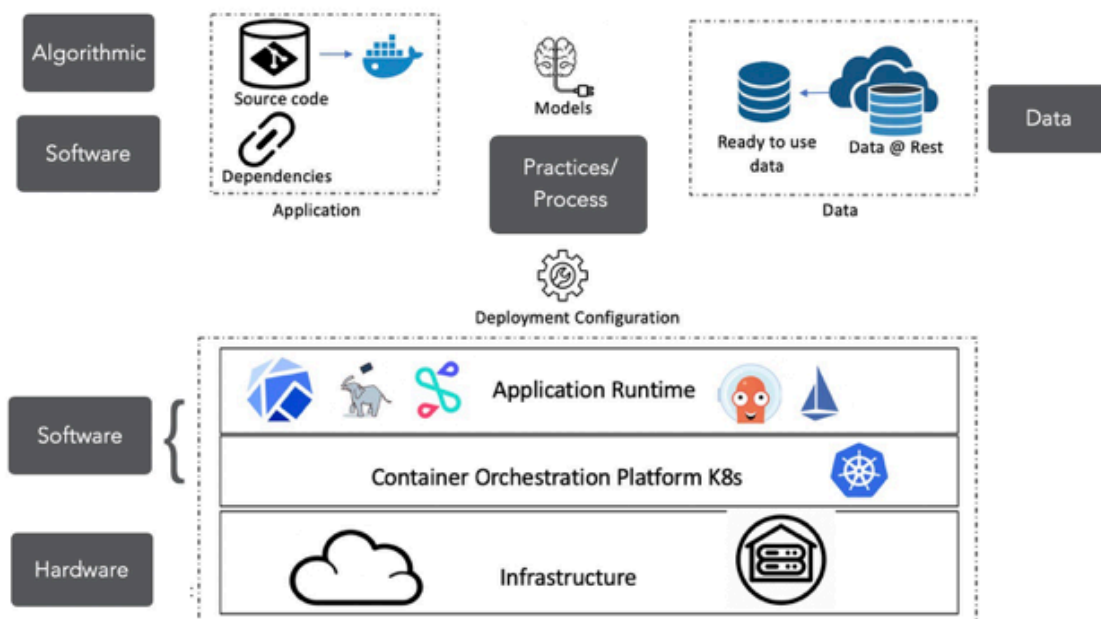


Figure 4: What to version control?

But first lets talk about creating the environment, infrastructure and versioning it.

## 1. Versioning environment

Using [gitops](#) environment and any changes associated with it can be version controlled. In this sample, we will be using [ArgoCD](#) to implement gitops workflow (figure 5) which will see our environment config moved to be alongside of code repository (figure 6).

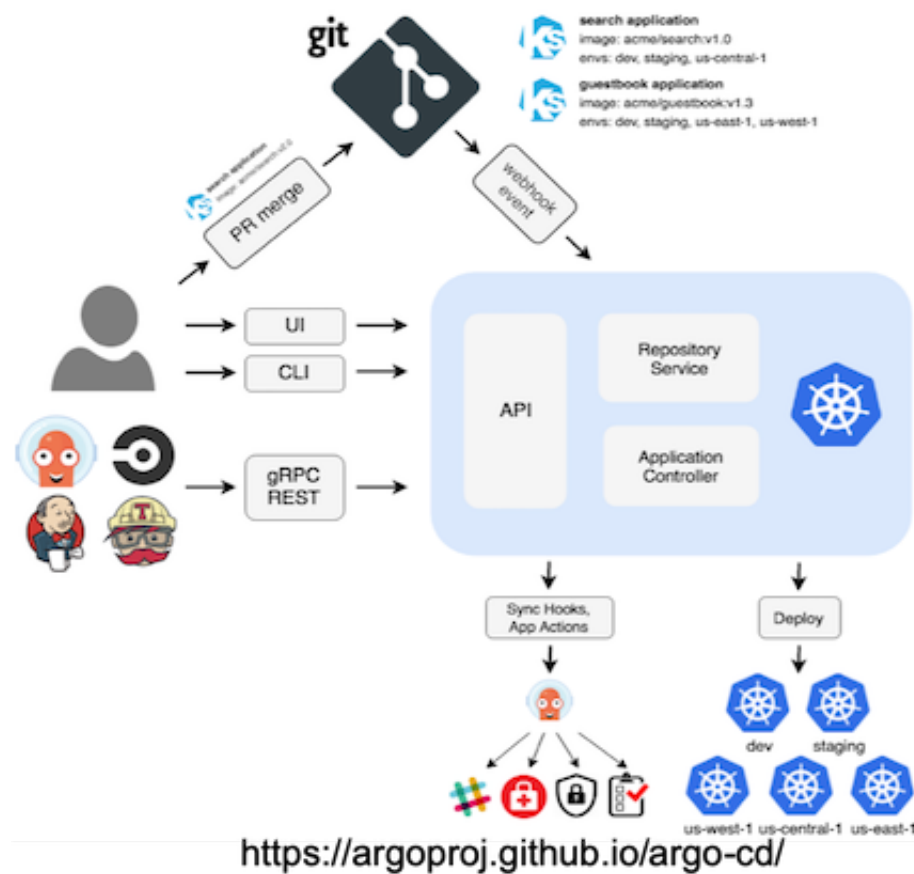


Figure 5: Gitops

This is achieved by defining [argo apps](#) which can be applied on BYO Kubernetes cluster (version 1.14.7) that has ArgoCD installed:

```
kubectl apply -f https://raw.githubusercontent.com/suneeta-mall/e2e-ml-on-k8s/master
```

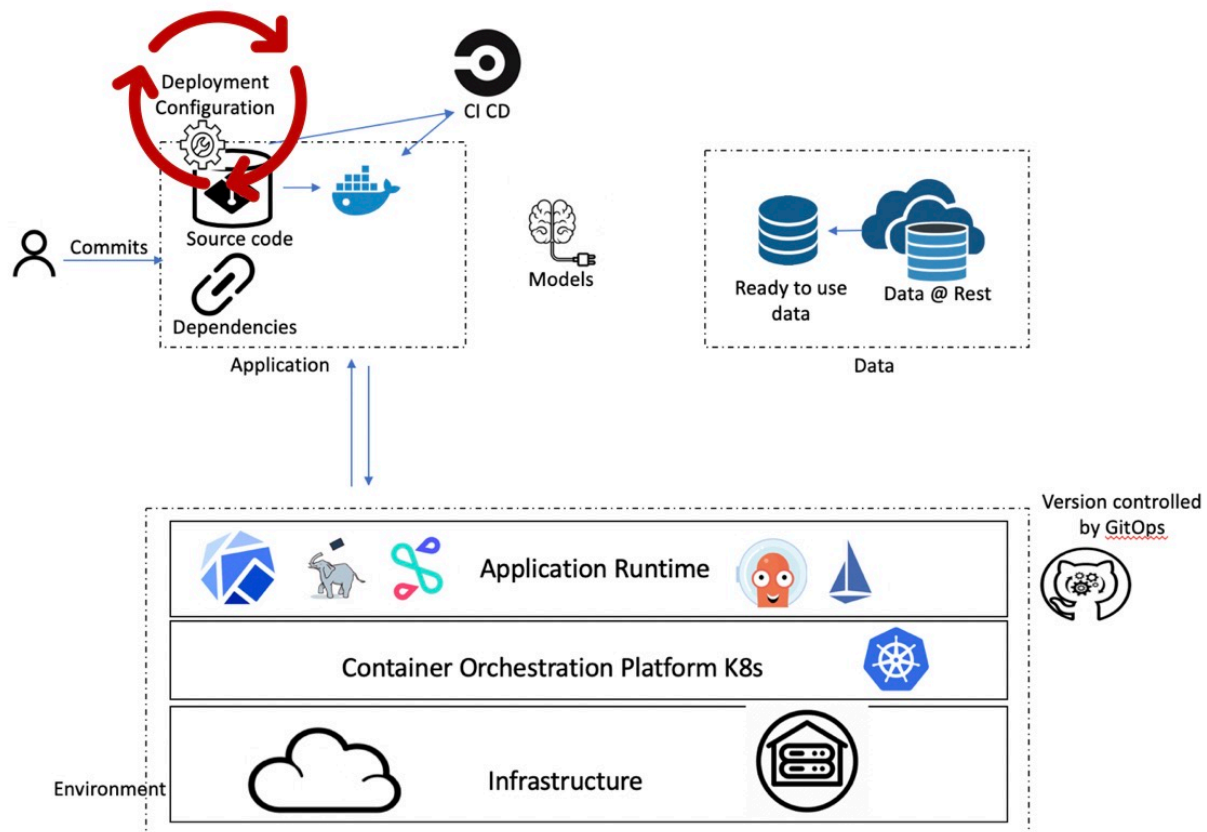


Figure 6: Gitops on environment config

Once the Argo apps are created, following softwares will be installed on the cluster:

- Kubernetes: 1.14.7 (tested on this version, in theory should work with other versions too!)
- ArgoCD: 1.2.3
- Kubeflow: 0.6.2

Kubeflow is ML toolkit designed to bring variety of ML related **Kubernetes** based softwares together.

- Seldon 0.4.1 (upgraded from packaged version on kubeflow 0.6.2)

Seldon is model serving software

- Pachyderm: 1.9.7

Pachyderm offers git like repository that can hold data even big data. It also offers automated repository capability that act on input and generate data thus holding versioned copy of this generated data. Together with these constructs it can be used to create pipeline DAG like processes with provenance across graph input, transformation spec, output

Any change on this configuration repository will then trigger a cluster update keeping environment in synced with versioned config.

## 2. Versioning data, process and artifacts

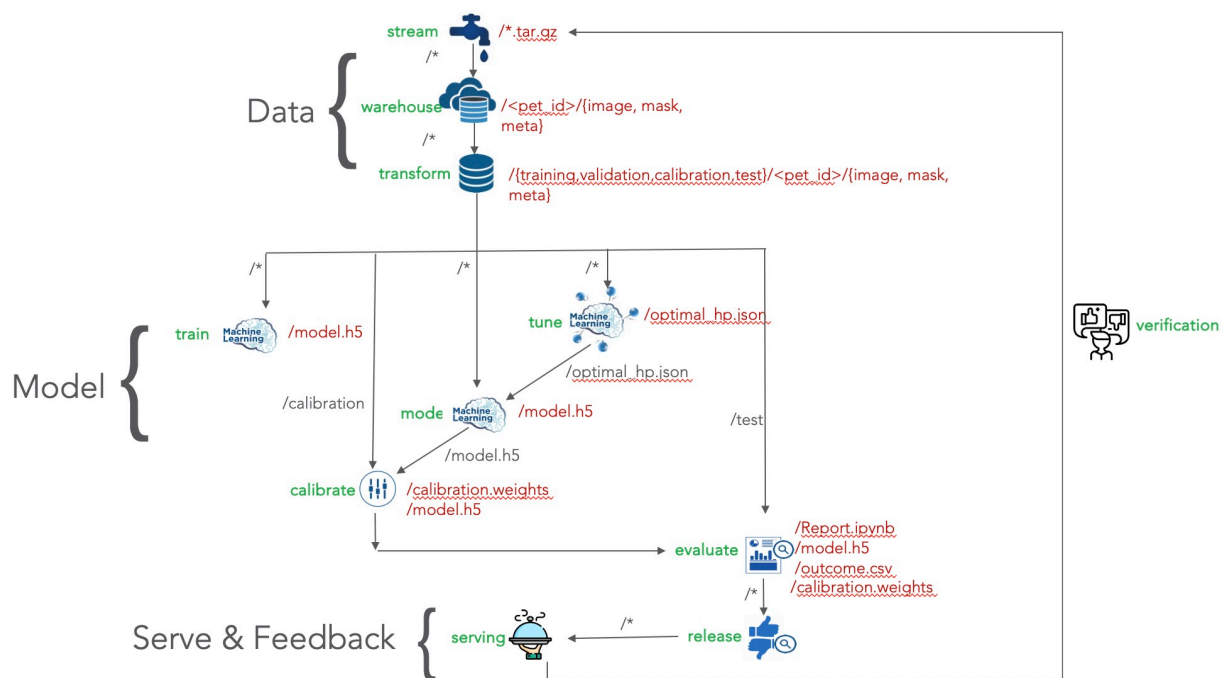


Figure 7: Artifact view of Machine Learning end to end system (shown in figure 2)

Pachyderm pipeline specification for end to end ML workflow capability shown in figure 2 is available [here](#). The generated artifacts/data as a result of this ML workflow is shown in figure 7 above. These artifacts and their association with other processes are also highlighted in figure 7.

```

---
pipeline:
  name: stream
transform:
  image: suneetamall/pykubectl:1.14.7
  cmd:
    - "/bin/bash"
  stdin:
    - "wget -O images.tar.gz https://www.robots.ox.ac.uk/~vgg/data/pets/data/images
      wget -O annotations.tar.gz https://www.robots.ox.ac.uk/~vgg/data/pets/data/ar
      tar -cvf data.tar.gz *.tar.gz && \
      cat data.tar.gz > /pfs/out && \
      while ;; do sleep 2073600; done"
spout:
  overwrite: true
---
input:
  pfs:
    glob: /
  
```

```

    repo: stream
pipeline:
  name: warehouse
transform:
  cmd:
    - "/bin/bash"
  image: suneetamall/e2e-ml-on-k8s:1
  stdin:
    - "python download_petset.py --input /pfs/stream/ --output /pfs/out"
datum_tries: 2
#standby: true
---
input:
  pfs:
    glob: "/"
    repo: warehouse
pipeline:
  name: transform
transform:
  cmd:
    - "/bin/bash"
  image: suneetamall/e2e-ml-on-k8s:1
  stdin:
    - "python dataset_gen.py --input /pfs/warehouse --output /pfs/out"
datum_tries: 2
#standby: true
---
input:
  pfs:
    glob: "/"
    repo: transform
pipeline:
  name: train
transform:
  cmd:
    - "/bin/bash"
  image: suneetamall/e2e-ml-on-k8s:1
  stdin:
    - "python train.py --input /pfs/transform --output /pfs/out --checkpoint_path /i
resource_requests:
  memory: 2G
# gpu:
#   type: nvidia.com/gpu
#   number: 1
datum_tries: 2
#standby: true
---
input:

```

```

    pfs:
      glob: "/"
      repo: transform
  pipeline:
    name: tune
  transform:
    cmd:
      - "/bin/bash"
    image: suneetamall/e2e-ml-on-k8s:1
    stdin:
      - "python tune.py --input /pfs/transform --output /pfs/out"
  resource_requests:
    memory: 4G
    cpu: 1
#   gpu:
#     type: nvidia.com/gpu
#     number: 1
  datum_tries: 2
#standby: true
---
input:
  cross:
    - pfs:
        glob: "/"
        repo: transform
    - pfs:
        glob: "/optimal_hp.json"
        repo: tune
  pipeline:
    name: model
  transform:
    cmd:
      - "/bin/bash"
    image: suneetamall/e2e-ml-on-k8s:1
    stdin:
      - "python train.py --input /pfs/transform --hyperparam_fn_path /pfs/tune/optimal
        --output /pfs/out --checkpoint_path /pfs/out/ckpts --tensorboard_path /pfs/out
      - "ln -s /pfs/tune/optimal_hp.json /pfs/out/optimal_hp.json"
  resource_requests:
    memory: 2G
#   gpu:
#     type: nvidia.com/gpu
#     number: 1
  datum_tries: 2
#standby: true
---
input:
  cross:

```



```

- pfs:
  glob: "/calibration"
  repo: transform
- pfs:
  glob: "/model.h5"
  repo: model
pipeline:
  name: calibrate
transform:
  cmd:
  - "/bin/bash"
  image: suneetamall/e2e-ml-on-k8s:1
  stdin:
  - "python calibrate.py --input /pfs/transform --model_weight /pfs/model/model.h5"
  - "ln -s /pfs/model/model.h5 /pfs/out/model.h5"
datum_tries: 2
#standby: true
---
input:
  cross:
  - pfs:
    glob: "/test"
    repo: transform
  - pfs:
    glob: "/"
    repo: calibrate
pipeline:
  name: evaluate
transform:
  cmd:
  - "/bin/bash"
  image: suneetamall/e2e-ml-on-k8s:1
  stdin:
  - "papermill evaluate.ipynb /pfs/out/Report.ipynb \
    -p model_weights /pfs/calibrate/model.h5 \
    -p calibration_weights /pfs/calibrate/calibration.weights \
    -p input_data_dir /pfs/transform \
    -p out_dir /pfs/out \
    -p hyperparameters /pfs/calibrate/optimal_hp.json"
  - "ln -s /pfs/calibrate/model.h5 /pfs/out/model.h5"
  - "ln -s /pfs/calibrate/calibration.weights /pfs/out/calibration.weights"
resource_requests:
  memory: 1G
datum_tries: 2
#standby: true
---
input:
  pfs:

```

```

    glob: "/"
    repo: evaluate
pipeline:
  name: release
transform:
  cmd:
    - "/bin/bash"
  image: suneetamall/e2e-ml-on-k8s:1
  stdin:
    - "python release.py --model_db evaluate --input /pfs/evaluate/evaluation_result"
pod_spec: '{"serviceAccount": "ml-user", "serviceAccountName": "ml-user"}'
datum_tries: 2
#standby: true
---
## Service https://docs.pachyderm.io/en/1/concepts/pipeline-concepts/pipeline/service
input:
  pfs:
    glob: "/"
    repo: model
pipeline:
  name: tensorboard
service:
  external_port: 30888
  internal_port: 6006
transform:
  cmd:
    - "/bin/bash"
  stdin:
    - tensorboard --logdir=/pfs/model/
  image: suneetamall/e2e-ml-on-k8s:1
---
```

This pipeline creates ML workflow, with artifact dependency shown in above figure 7, wherein full provenance across data, processes and outcome is maintained along with respective lineage.

This is the last post of the technical blog series, [Reproducibility in Machine Learning](#).

Realizing reproducible  
Machine Learning - with  
Tensorflow »

suneeta-mall © 2010-2019 , subscribe.

Powered by Jekyll & Polar